

Runtime Environment

Topics we will cover

- The procedure abstraction and linkage conventions
- Runtime storage convention
- Non-local data access (brief)

These issues are critical to high-performance code generation

Topics not covered :

- Garbage collection support
- Heap management
- Exception handling support
- Optimization for cache and TLB: covered next semester

Nevertheless, these issues are important for performance.

The Procedure Abstraction and Separate Compilation

Requires system-wide compact

- must involve architecture, OS, and compiler
- broad agreement on memory layout, protection requirements, calling sequence, error handling and reporting

Separate compilation

- *Compilation strategy that allows subsets (files, modules, or directories) of a single program to be compiled separately and then linked together*
- essential for building large systems
- keeps compile times reasonable
- requires independent procedures

Establishes the need for private context

- create a run-time "record" for each procedure invocation
- encapsulate run-time control and data information

Storage Management Conventions

1. Storage layout convention

- defines memory layout of *code* and *data* (static, stack, heap)
- mostly specified by operating system

2. Linkage convention

- protocol for passing values and program control at procedure call and return
- all code in a single program must follow common convention
- partly specified by processor and OS; rest is left to compiler


Related Tools

- **linker:**
 - resolve name references across procedures
 - static linking: absolute virtual addresses chosen at link-time
 - dynamic linking: some addresses chosen at load time
- **loader:** loads code and static data into memory

The linkage convention

The linkage convention ensures that procedures inherit a valid run-time environment and that they restore one for their parents.

- specifies steps in calling sequence and return sequence
- division of responsibility between caller and callee



At compile-time: generate code to implement linkage convention
At run-time: that code manipulates the stack frame and data areas


Activation Record or Stack Frame

At *run-time*, each procedure invocation has some associated storage. We call this its activation record. For most languages, activation records can live on the stack, and then they are also called stack frames.

Why is this useful?

Common components of a stack frame?

Example Stack Frame Layout



Assumptions

- Stack grows downwards in memory
- Dedicated frame pointer (FP) register
- Stack pointer (SP), if any, is separate

Also see Sparc V9 Stack Frame handout.

Procedure linkages

At a call

Caller:

- allocate basic frame
- store parameters
- store return address
- save *caller-saved* regs
- store self's FP
- set FP for child
- jump to child

Callee:

- save *callee-saved* regs, state
- extend frame (for locals)
- initialize locals
- fall through to code

At a return

Caller:


- copy return value
- deallocate basic frame
- restore *caller-saved* regs.

Callee:

- store return value
- restore *callee-saved* regs, state
- unextend frame
- restore parent's FP
- jump to return address

Run-time Storage Organization

Logical address-space



Code

- functions*: fixed size, statically allocated

Control stack

- holds frames at run-time: dynamic slice of the *activation tree*

Data

- initialized global data
- uninitialized global data
- statically allocated local data
- variable-size or dynamic local data
- dynamic non-local data

Mapping variables into memory

Binding \equiv mapping of a name to an attribute (e.g., storage location)

Compile time

- assign each name a class (*base address*) & offset (*location*)
- *Code generation*: generate instruction sequence for each access

Link time

- decide static memory layout
- resolve references across procedures (global variables and functions)

Load time

- load code and static data into memory
- resolve relocatable labels

Run time

- execute linkage code at each procedure invocation
- execute generated instruction sequences

Run-time Storage Organization

Each variable must be assigned a storage class

Global and “static” variables

- symbolic addresses compiled into code (*relocatable*)
- limited to fixed-size objects
- layout may be important for performance
- compiler enforces access rights (private vs. global)

Procedure-local variables

- Put in the stack frame *if* lifetime is limited (*see next slide*)
- Otherwise allocate on the heap
- *On stack*: May be allocated statically or dynamically
- *On stack*: May be fixed-size or variable size
- *On stack*: Implicit deallocation at procedure return

Dynamically Allocated (non-local) Variables

- returned pointers lead to non-local lifetimes
- explicit allocation
- explicit or implicit deallocation

Run-time Storage Organization

Do local variables go on stack?
 \Leftarrow Lifetime of memory is key

Downward exposure

- called procedure may reference my variables
- lexical or dynamic scoping

Upward exposure

- return a reference to my variables
- return a function that may reference my variables
- special case: *later instance* of current function or callee may reference my variables
 - e.g., `static` (C) or `save` (FORTRAN)
 - problem: value must be preserved across calls

Run-time Storage Organization

How should these variables be allocated? Assume all are local/formal.

```

1      QueueItem* q;           /* q? */
2      static int n;           /* n? */

```

```

1      QueueItem q = new QueueItem; // q? new object?

```

```

1      let m: Int <- a+b in self.n <- m -- m? self?

```

```

1      # let findFunc = fun x -> (* x? *)
2      let p = makePair (x+1) (x+2) in (* p? *)
3      (fun y -> (car p > y));; (* y? fun? *)

```

Run-time Storage Organization

Conditions for Stack Allocation

Allocate procedure-local data values in the stack frame if the locations are downwards-exposed only
 \Rightarrow values are not preserved across calls

Optimization

- **Escape analysis:** Decides whether a local variable or object can be used after procedure returns.
- Important for languages like Java, C# where all objects are heap-allocated by default.

Access to non-local data

Global & static variables

Compiler

- assign variable an offset (k) from some symbolic address (*label*)
- generate unique label for each name (*name mangling*)
- code to emit:

```
loadI    <label>    → r1
loadAI   r1, k,     → r2
```

Static linking

- Resolve all labels at link-time
Relocation
- Assigns each label an absolute virtual address

Dynamic linking

- Relocate some labels at run-time
- *Indirection table for external references*
- **Compiler:** generates extra load to get base address from table
- **Runtime linker:** fills addresses of labels into indirection table

Access to non-local data

Local Variables with Lexical Scoping

Languages without nested procedures :

- C, C++, Java, etc.
- Non-local names must be global or static
- Note: Heap-allocated data has no name: need some non-heap variable to access it (e.g., C pointer or Java reference)

Languages with nested procedures (e.g., Pascal) :

- view variables as $\langle level, offset \rangle$ pairs (compile-time)
- find pointer to appropriate activation record for *level*
- add *offset* to *level's* pointer
- more expensive to access than locals

The $\langle level, offset \rangle$ pair is called a Static Distance Coordinate

Access to non-local data

Two important problems arise

- How do we map a name into a $\langle level, offset \rangle$ pair?
- Given a $\langle level, offset \rangle$ pair, how to compute the address?

How do we map a name into a level & offset?

Use a *scoped symbol table*

(compile-time)

- look up a name, want its most recent declaration
- declaration may be at current level or any lower level
- offsets directly determined by stack frame layout

Given a level & offset, what's the address?

Two classic approaches

(run-time)

- access links or static links
- displays

See slides and Cooper-Torczon text if interested.

(run-time)

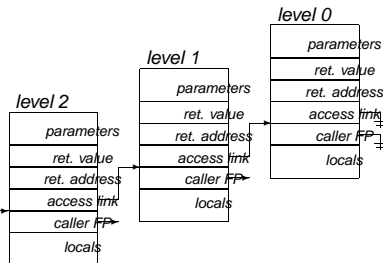
Access links: run-time addressability

The idea

- each frame contains a pointer to its lexical parent
- chain of indirection follows lexical nesting

The implementation

- creates a chain of frames by lexical ancestry
- can follow chain to find level k frame
- must maintain chain at call & return



Note: caller FP is not same as access link

Access links: run-time addressability

To find the value specified by $\langle l, o \rangle$

(assume $k =$ current procedure level = 2)

$\langle l, o \rangle$	Generated code
$\langle 2, 8 \rangle$	loadAI FP, 8 \Rightarrow r2
$\langle 1, 12 \rangle$	loadAI FP, -4 \Rightarrow r1 loadAI r1, 12 \Rightarrow r2
$\langle 0, 16 \rangle$	loadAI FP, -4 \Rightarrow r1 loadAI r1, -4 \Rightarrow r1 loadAI r1, 16 \Rightarrow r2

Access cost for non-locals varies with $k - l$ Note: $k < l$ cannot occur

Maintaining access links: (static links)

- calling level $k + 1$ procedure
 - pass my FP as access link
 - my backward chain will work for lower levels
- calling procedure at level $l < k$
 - find my link to level $l - 1$ and pass it
 - its access link will work for lower levels

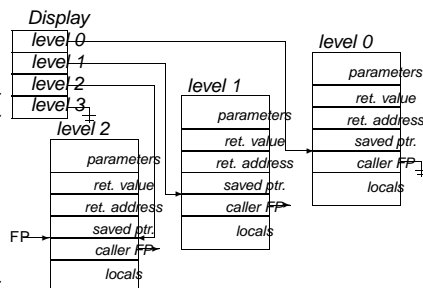
The display: run-time addressability

The idea

- replace linked list with global array indexed by level
- global array: holds currently active frame pointers

The implementation

- single, global table
- each call updates the table appropriately
- reference is one indirection through table



The display: run-time addressability

To find the value specified by $\langle l, o \rangle$

(assume $k = 2$)

$\langle l, o \rangle$	Generated code
$\langle 2, 8 \rangle$	loadAI FP, 8 \Rightarrow r2
$\langle 1, 12 \rangle$	loadI DISPLAY_BASE \Rightarrow r1 loadAI r1, 4 \Rightarrow r1 loadAI r1, 12 \Rightarrow r2
$\langle 0, 16 \rangle$	loadI DISPLAY_BASE \Rightarrow r1 loadAI r1, 8 \Rightarrow r1 loadAI r1, 16 \Rightarrow r2

Desired FP is at $DB + 4 \times l$

Access cost is constant & independent of $k - l$

Maintaining a global display (overallocate by 1 slot)

- on entry to procedure at level l
 - save level l display ptr in frame
 - push FP into level l display slot
- on return
 - restore the level l display slot from frame

Quick, simple, & foolproof

Minor issues (Displays & access links)

Improvements

- *Leaf procedures* (contain no calls)
 - with display, don't update it
 - statically allocate frame
- *Other procedures*
 - 1 call in loop \Rightarrow move frame manipulation out of loop
 - keep accessed FP elements in temporaries

Cost comparison

- *Display* *Major problem: threads*
 - load + store on CALL & RETURN
 - loadI + 2 loadAI's per access
 - *Access links*
 - store on CALL, nothing on RETURN
 - $(k - l + 1)$ loadAI's per access
-