


Instruction scheduling: The engineer's view

The problem

Given a code fragment for some target machine and the latencies for each individual instruction, reorder the instructions to minimize execution time

Conceptually, a scheduler looks like



Its task

- produce correct code (preserve flow of data)
- minimize wasted cycles (interlocks & stalls)
- avoid spilling registers (adding stores & loads)
- operate efficiently (reasonable compile time)

CS326

-1-

Instruction Scheduling

Instruction scheduling: The abstract view

Scheduling graph

To capture the important properties of the code, we build a scheduling graph, $G = (N, E, type, delay)$.

Each $n \in N$ is an instruction of $type(n)$ with $delay(n)$.

An edge $e = (n_1, n_2) \in E$ iff n_2 uses n_1 .

Definitions

A correct schedule S maps each $n \in N$ into a non-negative integer that represents its cycle number, and

1. $S(n) \geq 0$, for all $n \in N$
2. if $(n_1, n_2) \in E$, $S(n_1) + delay(n_1) \leq S(n_2)$
3. for each type t , there are no more instructions of type t in any cycle than the machine can issue

The *length* of a schedule S , denoted $L(S)$, is

$$L(S) = \max_{n \in N} (S(n) + delay(n))$$

The goal is to find the shortest possible correct schedule.

S is optimal if $L(S) \leq L(S_1)$, \forall schedules S_1 .

CS326

-3-

Instruction Scheduling

Example

$w \leftarrow w * 2 * x * y * z$

Assume:

- non-blocking load
- arguments can be reused next cycle

Cycles per op	
<i>fload</i>	5
<i>fstore</i>	5
<i>floadi</i>	1
<i>fshift</i>	1
<i>fadd</i>	1
<i>fmult</i>	2

<i>simple schedule</i>	<i>load aggressively</i>
1 <i>fload</i> $r_1 \leftarrow sp+@w$	1 <i>fload</i> $r_1 \leftarrow sp+@w$
2 <i>floadi</i> $r_2 \leftarrow 2$	2 <i>fload</i> $r_2 \leftarrow sp+@x$
6 <i>fmult</i> $r_1 \leftarrow r_1, r_2$	3 <i>fload</i> $r_3 \leftarrow sp+@y$
7 <i>fload</i> $r_2 \leftarrow sp+@x$	4 <i>fload</i> $r_4 \leftarrow sp+@z$
12 <i>fmult</i> $r_1 \leftarrow r_1, r_2$	5 <i>floadi</i> $r_5 \leftarrow 2$
13 <i>fload</i> $r_2 \leftarrow sp+@y$	6 <i>fmult</i> $r_1 \leftarrow r_1, r_5$
18 <i>fmult</i> $r_1 \leftarrow r_1, r_2$	8 <i>fmult</i> $r_1 \leftarrow r_1, r_2$
19 <i>fload</i> $r_2 \leftarrow sp+@z$	10 <i>fmult</i> $r_1 \leftarrow r_1, r_3$
24 <i>fmult</i> $r_1 \leftarrow r_1, r_2$	12 <i>fmult</i> $r_1 \leftarrow r_1, r_4$
26 <i>fstore</i> $sp+@2 \leftarrow r_1$	14 <i>fstore</i> $sp+@2 \leftarrow r_1$
33 r_1 available again	15 ...
2 registers, 33 cycles	19 r_1 available again
	5 registers, 19 cycles

Heuristics

(Proebsting)

- hiding latency of k requires $k + 1$ registers
- load aggressively, fill with operations

CS326

-2-

Instruction Scheduling

Instruction scheduling: What's so difficult?

Critical points

1. operands must be available (correctness)
 2. multiple ops can be *ready* (choice)
 3. moving ops can lengthen register lifetimes
 4. uses near definitions can shorten register lifetimes
 5. ops have multiple predecessors (start of block)
- Together, these issues make scheduling hard (NP-complete)

Simple case

- restricted to straight-line code
- single instruction per cycle
- consistent and predictable latencies

Even the simple case is NP-complete

CS326

-4-

Instruction Scheduling

1. “**Labelling**” algorithm [Sethi and Ullman]:
Optimal code assuming single-cycle loads
 - ignores load latency
 - assumes *expression tree* of entire basic-block (*not dag*)
 - combines instruction scheduling and register allocation
2. **DLS algorithm** [Proebsting and Fischer]:
Optimal code for a delayed-load architecture
 - fixed multi-cycle load latency
 - still assumes expression tree
 - Direct extension of Sethi-Ullman
3. **List-scheduling algorithm** [Gibbons and Muchnick]:
Optimal scheduling for pipelined multiple-issue processors
 - fixed multi-cycle load latency
 - takes linear IL as input (e.g., 3-address code)
 - linear IL \Rightarrow prior optimizations possible
 - linear IL \Rightarrow some register allocation may be done already

List scheduling algorithm: Engineering

Use clever data structures:

1. **ReadyL**: single ready list
2. $W[c]$, $0 \leq c < \text{MaxExecTime}$: worklists by cycle,
 $\text{MaxExecTime} \equiv \max_{n \in \text{DAG}} \text{ExecTime}(n)$
3. For each edge $e : n \rightarrow s$:
 $\text{Avail}(s,e) \equiv$ cycle in which value from node n is available to
node s (at the *start* of the cycle)

Heuristic function **ChooseInstr**(c , **ReadyL**, DAG):

- choose instruction from **ReadyL** to schedule in cycle c
- delete the instruction from **ReadyL**
- heuristics used here are key to good performance (later)

Simple idea

1. retain a ready list of instructions by cycle
2. repeat cycle-by-cycle until all instructions scheduled:
 - (a) choose an instruction and schedule it
 - (b) add successors to appropriate ready list

But “list scheduling” is really a class of algorithms that use different heuristics for step 2(a).

Input

- $DAG(N, E)$ for basic block
- $\text{ExecTime}(n) \equiv$ latency for each node (instruction) $n \in DAG$
(Can add specific latencies between pairs of instruction types to model more complex resource conflicts.)

Output

- **Start**(n) \equiv cycle in which instruction at node n begins execution

See example on slide 12.

List scheduling algorithm: details

I. Initialization:

1. for each instruction i in block
 - a. initialize **Avail**(i,e) appropriately *How?*
 - b. if i is ready, add it to **ReadyL** *Which i?*
2. $\forall 0 \leq c < \text{MaxExecTime}$, $W[c] \leftarrow \emptyset$
3. cycle $\leftarrow 1$

II. Repeat until all instructions are scheduled:

1. $i \leftarrow \text{ChooseInstr}(\text{cycle}, \text{ReadyL}, \text{DAG})$ *What if no such i?*
2. **Start**(i) = cycle
3. for each outgoing edge $e : i \rightarrow s$
 - a. **Avail**(s,e) $\leftarrow \text{cycle} + \text{ExecTime}(i)$
 - b. if **Avail**(s,e) has been initialized for all edges coming in to s :
 - i. $c \leftarrow \text{MAX}_s \text{Avail}(s,e)$
 - ii. $c \leftarrow c \bmod \text{MaxExecTime}$ *Why?*
 - iii. $W[c] \leftarrow W[c] \cup s$ *Why?*
4. cycle $\leftarrow \text{cycle} + 1$
5. **ReadyL** $\leftarrow \text{ReadyL} \cup W[\text{cycle} \bmod \text{MaxExecTime}]$
6. $W[\text{cycle} \bmod \text{MaxExecTime}] \leftarrow \emptyset$

ChooseInstr(c, ReadyL, DAG)

Most common priority scheme:

Give priority to instructions on the critical path.

- Critical paths \equiv the longest path through the scheduling graph
- use depth-first traversal to compute path lengths
- replace worklist with priority queue $(+\log_2(n))$

Good tie-breakers are important for robustness

- rank by longest path containing each node
→ *priority to longest paths of original DAG*
- rank by number of successors in the DAG
→ *priority to nodes used by many other nodes*
→ *exposes most #candidates*
- go depth first in schedule graph
→ *minimize register lifetimes*
- schedule last use as soon as possible
→ *frees up a register quickly*

Forward and Backward List Scheduling

There are many variations of list scheduling algorithms, but they break down into two classes:

Forward list scheduling

- start with available ops
- work forward
- ready \Rightarrow all ops available

Backward list scheduling

- start with no successors
- work backward
- ready \Rightarrow latency covers uses

Which one is better?

- No clear choice: depends on dependence patterns, latencies
- A few critical operations may determine the overall schedule
- Critical operations appear near leaves: forward usually better
- Critical operations appear near roots: backward usually better

An idea: Why not try both and pick the best one?

- Building DAG and preprocessing is the biggest cost
- Scheduling algorithms themselves are relatively cheap
- Can try both forward and backward scheduling, and even multiple alternatives for each

See "Efficient instruction scheduling for a pipelined architecture" by P.B. Gibbons and S.S. Muchnick, *Proceedings of the SIGPLAN 86 Symposium on Compiler Construction*, July 1986 (SIGPLAN Notices 21(7))

Goal: avoiding pipeline hazards

- **load** followed by use of *that* register
- **store** followed by any **load**

Choice heuristics


1. instruction interlocks with successors in the *dag*
→ *interlock early \Rightarrow more candidates to cover it*
2. largest number of successors
3. longest path to roots of the *dag*

Results

- eliminates "most" pipeline hazards
- in practice, reasonably good schedules
- $O(n^2)$ complexity versus $O(n^4)$ for competition

This paper became one of the classics of scheduling literature.

List scheduling example: go from SPEC



opcode	ldi	lsl	add	addi	cmp	st
latency	1	1	2	1	1	4

Forward Schedule				Backward Schedule			
	Int.	Int.	Mem.		Int.	Int.	Mem.
1.	ldi ₁	lsl ₁	---	1.	ldi ₁	---	---
2.	ldi ₂	ldi ₃	---	2.	add ₁	lsl ₁	---
3.	ldi ₄	add ₁	---	3.	add ₁	ldi ₃	---
4.	add ₂	add ₃	---	4.	add ₃	ldi ₂	st ₅
5.	add ₄	add ₅	st ₁	5.	add ₂	ldi ₁	st ₄
6.	cmp	---	st ₂	6.	add ₁	---	st ₃
7.	---	---	st ₃	7.	---	---	st ₂
8.	---	---	st ₄	8.	---	---	st ₁
9.	---	---	st ₅	9.	---	---	---
10.	---	---	---	10.	---	---	---
11.	---	---	---	11.	cmp	---	---
12.	---	---	---	12.	br	---	---
13.	br	---	---				

List scheduling for extended basic blocks

- Single basic blocks are usually too small for wide-issue processors
- An extended basic block (EBB) \equiv a series of blocks, $b_1 \dots b_n$, where b_1 has multiple predecessors in the CFG, but $b_2 \dots b_n$ each has only one predecessor
- Apply list scheduling algorithm to each EBB at a time in the graph
- Be careful when moving code “before” a branch:
 - Move SSA register ops that cause no exceptions
 - Use speculative loads before branch

Trace scheduling

- Important for VLIW and very wide issue machines
 - \Leftarrow need to keep many functional units busy
- **Trace** \equiv an arbitrary sequence of basic blocks executed consecutively at runtime
- Use runtime profiles to choose most frequently executed traces
- Use list scheduling to schedule instructions on a trace, then eliminate its blocks from the graph, and move to the next trace
- More complicated rules for introducing copies

Software pipelining

- Critical for scheduling small loops on wide-issue processors
- Begins by folding loop to create longer loop bodies
 - \rightarrow e.g., fold once to execute 2 iterations concurrently
- Body of loop (the *kernel*) executes second half of iteration i and first half of iteration $i + 1$
 - \Rightarrow allows compiler to overlap long operations of iteration $i + 1$ with uses of iteration i
- Prolog executes first half of iteration 1; epilog executes last half of iteration n
- Use list scheduling to schedule the kernel
- Generalizes to 2 or more iterations in kernel
- Important supporting transformations:
 - Loop unrolling
 - Register renaming
 - Variable renaming when unrolling loops (if not SSA)